

Jungi Lee

✉ jungi.lee@snu.ac.kr 📍 Seoul, Republic of Korea 🏠 jungi-lee.github.io

EDUCATION

Seoul National University, Seoul, Republic of Korea

03/2023 - Present

M.S./Ph.D. in Electrical and Computer Engineering

Computer Architecture and Systems Lab, advised by Prof. Jaewoong Sim.

Seoul National University, Seoul, Republic of Korea

03/2017 - 02/2023

B.S. in Electrical and Computer Engineering

GPA: 3.87/4.30, Major GPA: 3.98/4.30

PUBLICATIONS

[OSDI '24] **InfiniGen: Efficient Generative Inference of Large Language Models with Dynamic KV Cache Management**

Wonbeom Lee*, Jungi Lee*, Junghwan Seo, and Jaewoong Sim

Proceedings of the 18th USENIX Symposium on Operating Systems Design and Implementation (OSDI), 2024

[ISCA '24] **Tender: Accelerating Large Language Models via Tensor Decomposition and Runtime Requantization**

Jungi Lee*, Wonbeom Lee*, and Jaewoong Sim

Proceedings of the 51st ACM/IEEE International Symposium on Computer Architecture (ISCA), 2024

[ASPLOS '24] **GSCore: Efficient Radiance Field Rendering via Architectural Support for 3D Gaussian Splatting**

Junseo Lee, Seokwon Lee, Jungi Lee, Junyong Park, and Jaewoong Sim

Proceedings of the 2024 International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), 2024

[ISCA '23] **NeuRex: A Case for Neural Rendering Acceleration**

Junseo Lee, Kwanseok Choi, Jungi Lee, Seokwon Lee, Joonho Whangbo, and Jaewoong Sim

Proceedings of the 50th ACM/IEEE International Symposium on Computer Architecture (ISCA), 2023

PROJECTS

- **Accelerator system for large language models through algorithm-hardware co-design.**
 - Under low-bit inference, it achieves up to an average of $2.63\times$ speedup over other accelerators with higher accuracy.
- **Dynamic key-value cache management solution for efficient generative inference in large language models.**
 - Novel KV cache management framework that provides scalability under long-text generation while achieving up to $3.00\times$ speedup over other management methods.

TEACHING EXPERIENCE

Seoul National University, Seoul, Republic of Korea

Graduate Teaching Assistant

- ECE 315.A - Digital Systems Design and Experiments
- ECE 322 - Computer Organization

Fall 2023

Spring 2023

Led lab sessions, answered questions about the class material, and graded labs/exams (ECE 315.A, ECE 322).

Guide course project on building a DNN accelerator on FPGA (ECE 315.A).

SKILLS

- **Languages:** C/C++, CUDA, Python
- **Applications/Frameworks:** PyTorch, TVM, Intel Pin, \LaTeX